

CAPSTONE PROJECT ON:

TOPIC MODELING ON NEWS ARTICLES

Presented by:
Subhadip Ghosh



Journey Roadmap:

- Problem Statement
- Discussing our documents
- Dataset Inspection
- Text Preprocessing
- Vectorization of tokens
- Topic Modeling using LDA (Latent Dirichlet Allocation)
- Discussing tokens in each topic
- Discussing our LDA visualization chart
- Conclusion
- Plans to improve

Problem Statement:

We are provided with a collection of articles corresponding to BBC News and the major topics/themes they have covered from 2004-2005


Objective of our model:

The objective of our model is to aggregate all these articles in one data frame and then identify these major topics/themes by clustering the appropriate terms that correctly correspond to the given themes.

Discussing our documents:



Business



Politics




Entertainment



Sports



Technology



Final dataset

Data Inspection:

- ▶ We have got 2225 documents in total
- ▶ We only have one column which is 'Articles' corresponding to the news reports of BBC News
- ▶ After aggregating all the documents, we resampled the dataset to mix up the articles
- ▶ By visually inspecting the first few documents of our dataset, we found a lot of unwanted characters like '/n, &, %, #'
- ▶ We need to remove all these special characters and all the stop words to work on only those terms which will be crucial to our topic-modelling algorithm.

Text Pre-processing

Tokenization

Removing
Stop-Words

Lemmatization

Checking
document
complexity

Tokenization:

- ▶ Tokenization is the process of turning the entire corpus into several 'tokens' or tiny pieces of terms which best describes the document when aggregated
- ▶ We used the split() method to tokenize our corpus and removed any unwanted characters and punctuations in our dataset
- ▶ We filtered out the crucial tokens by rejecting all the stopwords (is, and, the, are, they, etc...) by using the nltk library
- ▶ We used Lemmatization to convert these tokens into their base form in accordance to the English dictionary and these modified terms are called 'lemma'.

Example:


'neighbouring, neighbourhood, neighbours' turns into 'neighbour'

- ▶ Used Text blob to extract only the noun phrases from our corpus

Checking document complexity:

Complexity can be calculated for each document by:

- ▶ Calculating the number of characters in each document
- ▶ We will be looking into the distribution of character counts in our entire corpus



Vectorization:



Term vectorization



Checking the frequency of each token




Updating Stop-words and vectorizing using TF-IDF

Term-Vectorization:

- ▶ Unfortunately, our machine doesn't understand raw text data. These terms need to be converted into vectors for our machine to understand
- ▶ Term vectorization is the process of converting each and every token from the vocabulary into a corresponding vector of real numbers
- ▶ We will be using Count-Vectorizer to check the terms that occur most frequently in our corpus
- ▶ We will use the term frequency to update the stop-words and then apply TF-IDF for vectorization of the tokens

Most frequent terms:



```
[('mr', 2874),  
 ('year', 2198),  
 ('new', 1793),  
 ('people', 1252),  
 ('time', 1192),  
 ('game', 1170),  
 ('world', 996),  
 ('government', 954),  
 ('uk', 938),  
 ('film', 936),  
 ('company', 918),  
 ('bn', 871),  
 ('firm', 848),  
 ('service', 785),  
 ('market', 736),  
 ('number', 718),  
 ('labour', 693),  
 ('music', 678),  
 ('player', 672),
```

Updating Stopwords:

- ▶ In the previous section we checked the most frequently occurring words in our entire corpus
- ▶ The top three words “new, year, mr” won’t be of much help to us in topic modelling as these three words doesn’t correspond to any such major theme/topic that have been covered by BBC
- ▶ We have excluded all the words that have a frequency of more than 1500 and less than 4
- ▶ We added some more tokens in the stopwords list such as ‘abc, eu, th, aaa...etc’ as these tokens will also be of no help to us
- ▶ We will now use Tfidf-Vectorizer with these updated stopwords

Using TF-IDF Vectorizer:

Parameters used in TF-IDF Vectorizer:

- ▶ **ngram_range:** tells the model the range of word lengths it needs to consider while vectorization. We have used (1,2), meaning maximum two words can be collectively selected by the model
- ▶ **min_df:** Minimum document frequency. We have used 0.003, meaning any term present in less than 0.3 percent of the documents will not be considered
- ▶ **max_df:** Maximum document frequency. Used 0.60, meaning any term present in greater than 60 percent of the documents will not be considered
- ▶ After vectorizing using TF-IDF, we have a sparse matrix having shape (2215, 5647), meaning we have 2215 documents and 5647 tokens in the form of vectors

Fitting our Model

Libraries we
will be using

Fitting our
model using
LDA

Checking out
the topics

Discussing
LDA
visualisation
chart

Libraries we will be using and why:


- ▶ Importing pyLDAvis from sklearn
- ▶ Importing T-SNE (t-distributed stochastic neighbor embedding)
- ▶ Importing Latent Dirichlet Allocation

Topic Modeling using LDA


Selected Topic:

Slide to adjust relevance metric:(2)
 $\lambda = 0.59$

Intertopic Distance Map (via multidimensional scaling)



Top-30 Most Relevant Terms for Topic 4 (17.2% of tokens)




1. $\text{saliency}(\text{term } w) = \text{frequency}(w) * [\sum_t p(t | w) * \log(p(t | w)/p(t))]$ for topics t ; see Chuang et. al (2012)

2. $\text{relevance}(\text{term } w | \text{topic } t) = \lambda * p(w | t) + (1 - \lambda) * p(w | t)/p(w)$; see Sievert & Shirley (2014)


Terms in each topics:

Political


Agendas



Entertainment




Technology




Terms in each topic(contd.)

Sports



Economy



Discussing the LDA visualisation chart:

- ▶ The white bars define the overall term frequency of a particular token
- ▶ The red bars define the estimated term frequency of a particular token
- ▶ In each topic, the top 30 most relevant tokens have been used for visualization
- ▶ The relevance metric (lambda) when equal to 1, the terms are arranged according to their probability in a particular topic. Lambda when equal to zero, the tokens get arranged in the decreasing order of their respective lifts.

Lift: ratio of a term's probability to it's marginal probability in the entire corpus

- ▶ We have used T-SNE as a dimensionality reduction technique to visualise the clusters form in a 2-dimensional feature space

Conclusions:

- ▶ Out of the five major themes/topics that had been provided to us, our LDA (Latent Dirichlet Allocation) has correctly clustered most of the terms corresponding to a topic
- ▶ Business and Political articles constitute about 56% of the overall tokens and the rest three topics correspond to the remaining 44%
- ▶ Our LDA model works best at relevance level 0.6, correctly classifying almost 85% of the tokens in their respective clusters
- ▶ Adjusted the hyper-parameters `doc_topic_prior` and `topic_word_prior`, but for 5 topics, both corresponding to None gave the best results

Plans to improve:

- ▶ Most of the tokens from articles “Sports” have been removed either in the pre-processing stages or while updating the stop-words
- ▶ We could have use different parameters for TF-IDF for capturing more number of tokens by reducing the min_df even more. But due to higher computational powers required for a machine to generate a sparse matrix with more than 20,000 tokens, we were unable to test that out.